

CS 217 Data Management and Information Processing

Data Table and Pandas

Comma Separated Values (CSV)

- ▶ CSV is a simple text format for storing tabular data (spreadsheets)
- ▶ Each row is represented on one line of text
- ▶ Columns are separated by commas
- ▶ Values can be enclosed in double quotes ("...") if necessary
 - ▶ For example, if value includes comma or newline characters
 - ▶ Double quotes within a text value must be “escaped” by using two double quotes
- ▶ Values can be empty by having nothing between the commas

NBA_player_of_the_week.csv viewed in Excel

	A	B	C	D	E	F	G	H	I	J
1	PlayerID	TeamID	PositionID	First Name	Last Name	Seasons in L	Height	Weight	Age	
2	1	20	7	Micheal	Richardson	6	77	189	29	
3	2	14	9	Derek	Smith	2	78	205	23	
4	3	9	2	Calvin	Natt	5	79	220	28	
5	4	15	1	Kareem	Abdul-Jabbar	15	80	225	37	
6	5	2	8	Larry	Bird	5	81	220	28	
7	6	32	9	Darrell	Griffith	4	82	190	26	
8	7	11	7	Sleepy	Floyd	2	83	170	24	
9	8	8	8	Mark	Aguirre	3	84	232	25	
10	9	15	7	Magic	Johnson	5	85	255	25	
11	10	1	8	Dominique	Wilkins	2	86	200	25	
12	11	33	6	Tom	McMillen	9	87	215	32	
13	12	6	9	Michael	Jordan	0	88	215	22	
14	13	7	4	World	Free	9	89	185	31	
15	14	10	7	Isiah	Thomas	3	90	180	23	
16	15	18	6	Terry	Cummings	2	92	220	23	
17	16	6	6	Orlando	Woolridge	3	94	215	25	
18	17	30	1	Jack	Sikma	7	95	230	29	
19	18	22	8	Bernard	King	7	96	205	28	
20	19	25	1	Moses	Malone	8	97	215	29	
21	20	9	8	Alex	English	8	98	190	31	
22	21	26	6	Larry	Nance	3	99	205	26	
23	22	13	1	Herb	Williams	4	101	242	28	
24	23	25	6	Charles	Barkley	1	102	252	23	
25	24	32	8	Adrian	Dantley	9	85	208	30	
26	25	18	9	Sidney	Moncrief	6	89	180	28	

How to Process Tabular Data?

- ▶ Efficiency
 - ▶ The code should run quickly
- ▶ Easy to program
 - ▶ Should not take much effort to express our query
- ▶ Portable
 - ▶ The analysis can be quickly hooked up with other code blocks

Pandas

- ▶ One of the most popular library that data scientists use
- ▶ Created by Wes McKinney in 2008, now maintained by Jeff Reback and many others.
 - ▶ Author of one of the textbooks: Python for Data Analysis
- ▶ Powerful and productive Python data analysis and Management Library
- ▶ Its an open source product.
 - ▶ Free to use and free to modify



Overview

- ▶ Python Library to provide data analysis features similar to: R, MATLAB, SAS
- ▶ Rich data structures and functions to make working with data structure fast, easy and expressive.
- ▶ It is built on top of NumPy
- ▶ Key components provided by Pandas:
 - ▶ Series
 - ▶ DataFrame

**Might be the most frequently used tool
after taking this course!**

Pandas: Essential Concepts

- ▶ A **Series** is a named Python list (one-entry dict with list as value):

```
{ 'grades': [50,90,100,45] }
```

- ▶ A **DataFrame** is a collection of Series (dict-like container for series):

```
{'names': ['bob', 'ken', 'art', 'joe'],  
  'grades': [50,90,100,45]  
}
```

Series

- ▶ One dimensional array-like object
- ▶ It contains array of data (of any NumPy data type) with associated indexes. (Indexes can be strings or integers or other data types.)
- ▶ By default , the series will get indexing from 0 to N where $N = \text{size} - 1$

```
obj = Series([4, 7, -5, 3])  
obj
```

Output:

```
0    4  
1    7  
2   -5  
3    3  
dtype: int64
```

Index

Data

Series: Create

```
from pandas import Series
```

```
data = [1,2,3,numpy.nan,5,6]    # nan == Not a Number  
unindexed = Series(data)
```

```
indices = ['a', 'b', 'c', 'd', 'e', 'f']  
indexed = Series(data, index=indices)
```

```
data_dict = {'a' : 1, 'b' : 2, 'c' : 3}  
indexed = Series(data_dict)
```

Series: Accessing Elements

```
obj2 = Series([4, 7, -5, 3], \
              index=['d', 'b', 'a', 'c'])
```

obj2

Output:

```
d 4
b 7
a -5
c 3
dtype: int64
```

obj2.index

Output: Index(['d', 'b', 'a', 'c'],
dtype='object')

obj2.values

Output: array([4, 7, -5, 3], dtype=int64)

obj2['a']

Output: -5

obj2.a

Output: -5

obj2['d']=10

obj2[['d', 'c', 'a']]

Output:

```
d 10
c 3
a -5
dtype: int64
```

obj2[:2]

Output:

```
d 10
b 7
dtype: int64
```

Series - array/dict operations

- ▶ numpy array operations can also be applied, which will preserve the index-value link

```
obj2[obj2>0]
```

```
Output:
```

```
d 10
```

```
b 7
```

```
c 3
```

```
dtype: int64
```

```
obj2**2
```

```
Output:
```

```
d 100
```

```
b 49
```

```
a 25
```

```
c 9
```

```
dtype: int64
```

- ▶ Can be constructed from a dict directly.

```
obj3 = Series({'a': 10,  
'b': 5, 'c': 30})
```

```
obj3
```

```
Output:
```

```
a 10
```

```
b 5
```

```
c 30
```

```
dtype: int64
```

DataFrame

- ▶ A DataFrame is a tabular data structure comprised of rows and columns, akin to a spreadsheet or database table.
- ▶ It can be treated as an ordered collection of columns
 - ▶ Each column can be a different data type
 - ▶ Have both row and column indices

```
data = {'state': ['Ohio', 'Ohio',  
                'Ohio', 'Nevada', 'Nevada'], 'year':  
        [2000, 2001, 2002, 2001, 2002], 'pop':  
        [1.5, 1.7, 3.6, 2.4, 2.9]}  
frame = DataFrame(data)  
frame
```

Output:

	pop	state	year
0	1.5	Ohio	2000
1	1.7	Ohio	2001
2	3.6	Ohio	2002
3	2.4	Nevada	2001
4	2.9	Nevada	2002

DataFrame: Create

```
from pandas import DataFrame
```

```
data_dict = {'col1' : [1, 2, 3, 4],  
             'col2' : [10, 20, 30, 40]}
```

```
indices = ['a', 'b', 'c', 'd']
```

```
df = DataFrame(data_dict, index = indices)
```

df

Output:

	col1	col2
a	1	10
b	2	20
c	3	30
d	4	40

```
df2=DataFrame.from_items( [('col1', [1, 2, 3]),  
                           ('col2', [4, 5, 6])])
```

df2

Output:

	col1	col2
0	1	4
1	2	5
2	3	6

DataFrame: Create

```
pop = {'Nevada': {2001: 2.9, 2002: 2.9}, 'Ohio':  
{2002: 3.6, 2001: 1.7, 2000: 1.5}}
```

```
frame3 = DataFrame(pop)
```

```
frame3
```

Output:

	Nevada	Ohio
2000	NaN	1.5
2001	2.9	1.7
2002	2.9	3.6

DataFrame: index, columns, values

frame3.index

Output:

```
Int64Index([2000, 2001, 2002], dtype='int64')
```

frame3.columns

Output:

```
Index(['Nevada', 'Ohio'], dtype='object')
```

frame3.values

Output:

```
array([[ nan, 1.5],
       [ 2.9, 1.7],
       [ 2.9, 3.6]])
```

frame3

Output:

	Nevada	Ohio
2000	NaN	1.5
2001	2.9	1.7
2002	2.9	3.6

```
frame3.index.name = 'year'
frame3.columns.name='state'
frame3
```

Output:

```
state Nevada Ohio
year
2000 NaN 1.5
2001 2.9 1.7
2002 2.9 3.6
```

DataFrame: Retrieving a Column

- ▶ A column in a DataFrame can be retrieved as a Series by dict-like notation or as attribute
- ▶ Series index and name have been kept/set appropriately

```
frame['state']
```

Output:

0 Ohio

1 Ohio

2 Ohio

3 Nevada

4 Nevada

Name: state, dtype: object

```
type(frame['state'])
```

Output: pandas.core.series.Series

```
frame.state
```

Output:

0 Ohio

1 Ohio

2 Ohio

3 Nevada

4 Nevada

Name: state, dtype: object

frame

Output:

	pop	state	year
0	1.5	Ohio	2000
1	1.7	Ohio	2001
2	3.6	Ohio	2002
3	2.4	Nevada	2001
4	2.9	Nevada	2002

DataFrame: Getting Rows

- loc for using indexes and iloc for using positions

```
frame2
```

```
Output:
```

	year	state	pop	debt
A	2000	Ohio	1.5	NaN
B	2001	Ohio	1.7	NaN
C	2002	Ohio	3.6	NaN
D	2001	Nevada	2.4	NaN
E	2002	Nevada	2.9	NaN

```
frame2.loc['A']
```

```
Output:
```

```
year 2000
state Ohio
pop 1.5
debt NaN
Name: A, dtype:
object
```

```
type(frame2.loc['A'])
```

```
Output:
```

```
pandas.core.series.Series
```

```
frame2.loc[['A', 'B']]
```

```
Output:
```

	year	state	pop	debt
A	2000	Ohio	1.5	NaN
B	2001	Ohio	1.7	NaN

```
type(frame2.loc[['A', 'B']])
```

```
Output:
```

```
pandas.core.frame.DataFrame
```

More on DataFrame indexing

data

Output:

```
array([[0, 1, 2],  
       [3, 4, 5],  
       [6, 7, 8]])
```

```
frame = DataFrame(data,  
index=['r1', 'r2', 'r3'],  
columns=['c1', 'c2', 'c3'])
```

frame

Output:

	c1	c2	c3
r1	0	1	2
r2	3	4	5
r3	6	7	8

frame['c1']

Output:

```
r1 0  
r2 3  
r3 6  
Name: c1, dtype: int64
```

frame.loc['r1']

Output:

```
c1 0  
c2 1  
c3 2  
Name: r1, dtype:  
int64
```

frame['c1']['r1']

Output: 0

frame[['c1', 'c3']]

Output:

```
   c1 c3  
r1 0  2  
r2 3  5  
r3 6  8
```

frame.loc[['r1','r3']]

Output:

```
   c1 c2 c3  
r1 0  1  2  
r3 6  7  8
```

frame.iloc[:2]

Output:

```
   c1 c2 c3  
r1 0  1  2  
r2 3  4  5
```

More on DataFrame indexing - 2

frame

Output:

	c1	c2	c3
r1	0	1	2
r2	3	4	5
r3	6	7	8

frame[frame['c1']>0]

Output:

	c1	c2	c3
r2	3	4	5
r3	6	7	8

frame['c1']>0

Output:

r1 False
r2 True
r3 True

Name: c1, dtype: bool

frame < 3

Output:

	c1	c2	c3
r1	True	True	True
r2	False	False	False
r3	False	False	False

frame[frame<3] = 3

frame

Output:

	c1	c2	c3
r1	3	3	3
r2	3	4	5
r3	6	7	8

DataFrame - modifying columns

frame2

Output:

	year	state	pop	debt
A	2000	Ohio	1.5	NaN
B	2001	Ohio	1.7	NaN
C	2002	Ohio	3.6	NaN
D	2001	Nevada	2.4	NaN
E	2002	Nevada	2.9	NaN

Rows or individual elements can be modified similarly. Using loc or iloc.

```
val = Series([10, 10, 10],  
index = ['A', 'C', 'D'])
```

```
frame2['debt'] = 0
```

frame2

Output:

	year	state	pop	debt
A	2000	Ohio	1.5	0
B	2001	Ohio	1.7	0
C	2002	Ohio	3.6	0
D	2001	Nevada	2.4	0
E	2002	Nevada	2.9	0

```
frame2['debt'] = range(5)
```

frame2

Output:

	year	state	pop	debt
A	2000	Ohio	1.5	0
B	2001	Ohio	1.7	1
C	2002	Ohio	3.6	2
D	2001	Nevada	2.4	3
E	2002	Nevada	2.9	4

```
frame2['debt'] = val
```

frame2

Output:

	year	state	pop	debt
A	2000	Ohio	1.5	10.0
B	2001	Ohio	1.7	NaN
C	2002	Ohio	3.6	10.0
D	2001	Nevada	2.4	10.0
E	2002	Nevada	2.9	NaN

Removing rows/columns

```
frame.drop(['r1'])
```

Output:

	c1	c2	c3
r2	3	4	5
r3	6	7	8

```
frame.drop(['r1','r3'])
```

Output:

	c1	c2	c3
r2	3	4	5

```
frame.drop(['c1'], axis=1)
```

Output:

	c2	c3
r1	1	2
r2	4	5
r3	7	8

frame

Output:

	c1	c2	c3
r1	0	1	2
r2	3	4	5
r3	6	7	8

frame

Output:

	c1	c2	c3
r1	0	1	2
r2	3	4	5
r3	6	7	8

Does not change the old frame,
it returns a new frame

Function application and mapping

- ▶ `DataFrame.applymap(f)` applies `f` to every entry
- ▶ `DataFrame.apply(f)` applies `f` to every column or row

frame

Output:

	c1	c2	c3
r1	0	1	2
r2	3	4	5
r3	6	7	8

```
def square(x): return x**2
```

```
frame.applymap(square)
```

Output:

	c1	c2	c3
r1	0	1	4
r2	9	16	25
r3	36	49	64

```
def max_minus_min(x): return max(x)-min(x)
```

```
frame.apply(max_minus_min)
```

Output:

```
c1 6
c2 6
c3 6
dtype: int64
```

```
frame.apply(max_minus_min, axis=1)
```

Output:

```
r1 2
r2 2
r3 2
dtype: int64
```

Other DataFrame functions

- ▶ `head()` First few lines
- ▶ `tail(5)` Last 5 lines
- ▶ `mean()`
 - ▶ `Mean(axis=0, skipna=True)`
- ▶ `sum()`
- ▶ `describe()`: return summary statistics of each column
 - ▶ for numeric data: mean, std, max, min, 25%, 50%, 75%, etc.
 - ▶ For non-numeric data: count, uniq, most-frequent item, etc.

DataFrame: I/O

```
df = pandas.read_csv('data.csv')  
df.to_csv('data.csv')
```

```
df = pandas.read_excel('data.xlsx', 'Sheet1', index_col=None,  
na_values=['NA'])  
df.to_excel('data.xlsx', sheet_name='Sheet1')
```


Quiz

What is the VALUE and TYPE of each of the following?

1. `df['Quarter']`
2. `df[['Quarter']]`
3. `df[df['Quarter']=='Q1']`
4. `df[df['Sold'] < 110]`

► `df`

	Quarter	Sold
0	Q1	100
1	Q2	120
2	Q3	90
3	Q4	150

Quiz Answer

What is the VALUE and TYPE of each of the following?

1. `df['Quarter']`

```
>>> df['Quarter']
0    Q1
1    Q2
2    Q3
3    Q4
```

Series

2. `df[['Quarter']]`

```
>>> df[['Quarter']]
  Quarter
0      Q1
1      Q2
2      Q3
3      Q4
```

Dataframe

3. `df[df['Quarter']=='Q1']`

```
>>> df[df['Quarter']=='Q1']
  Quarter  Sold
0      Q1   100
```

Dataframe

4. `df[df['Sold'] < 110]`


```
>>> df[df['Sold'] < 110]
  Quarter  Sold
0      Q1   100
2      Q3    90
```

Dataframe

▶ df

	Quarter	Sold
0	Q1	100
1	Q2	120
2	Q3	90
3	Q4	150

Data Organization Spectrum



Starting from next
lecture!