

The background features abstract, overlapping green geometric shapes, primarily triangles and polygons, in various shades of green, creating a modern and dynamic visual effect.

# CS 217 Data Management and Information Processing

## OUTER JOINS and CROSS JOINS

# Last Week: GROUP BY and JOINS

- ▶ Default type of JOIN is the **INNER JOIN**
- ▶ Combines rows from two tables using a **join predicate**, which usually specifies that two columns must be equal.
- ▶ Multiple JOINS can be combined
- ▶ Refer to columns as *table.column*
- ▶ Can use AS to give a table an alias for use in the statement
  - ▶ Do this when joining a table two or more times, to distinguish each copy of the table.

# NATURAL JOIN

- ▶ A shorthand notation to make some JOINS shorter to express.
- ▶ NATURAL JOIN matches rows using whatever columns have identical names.

For example:

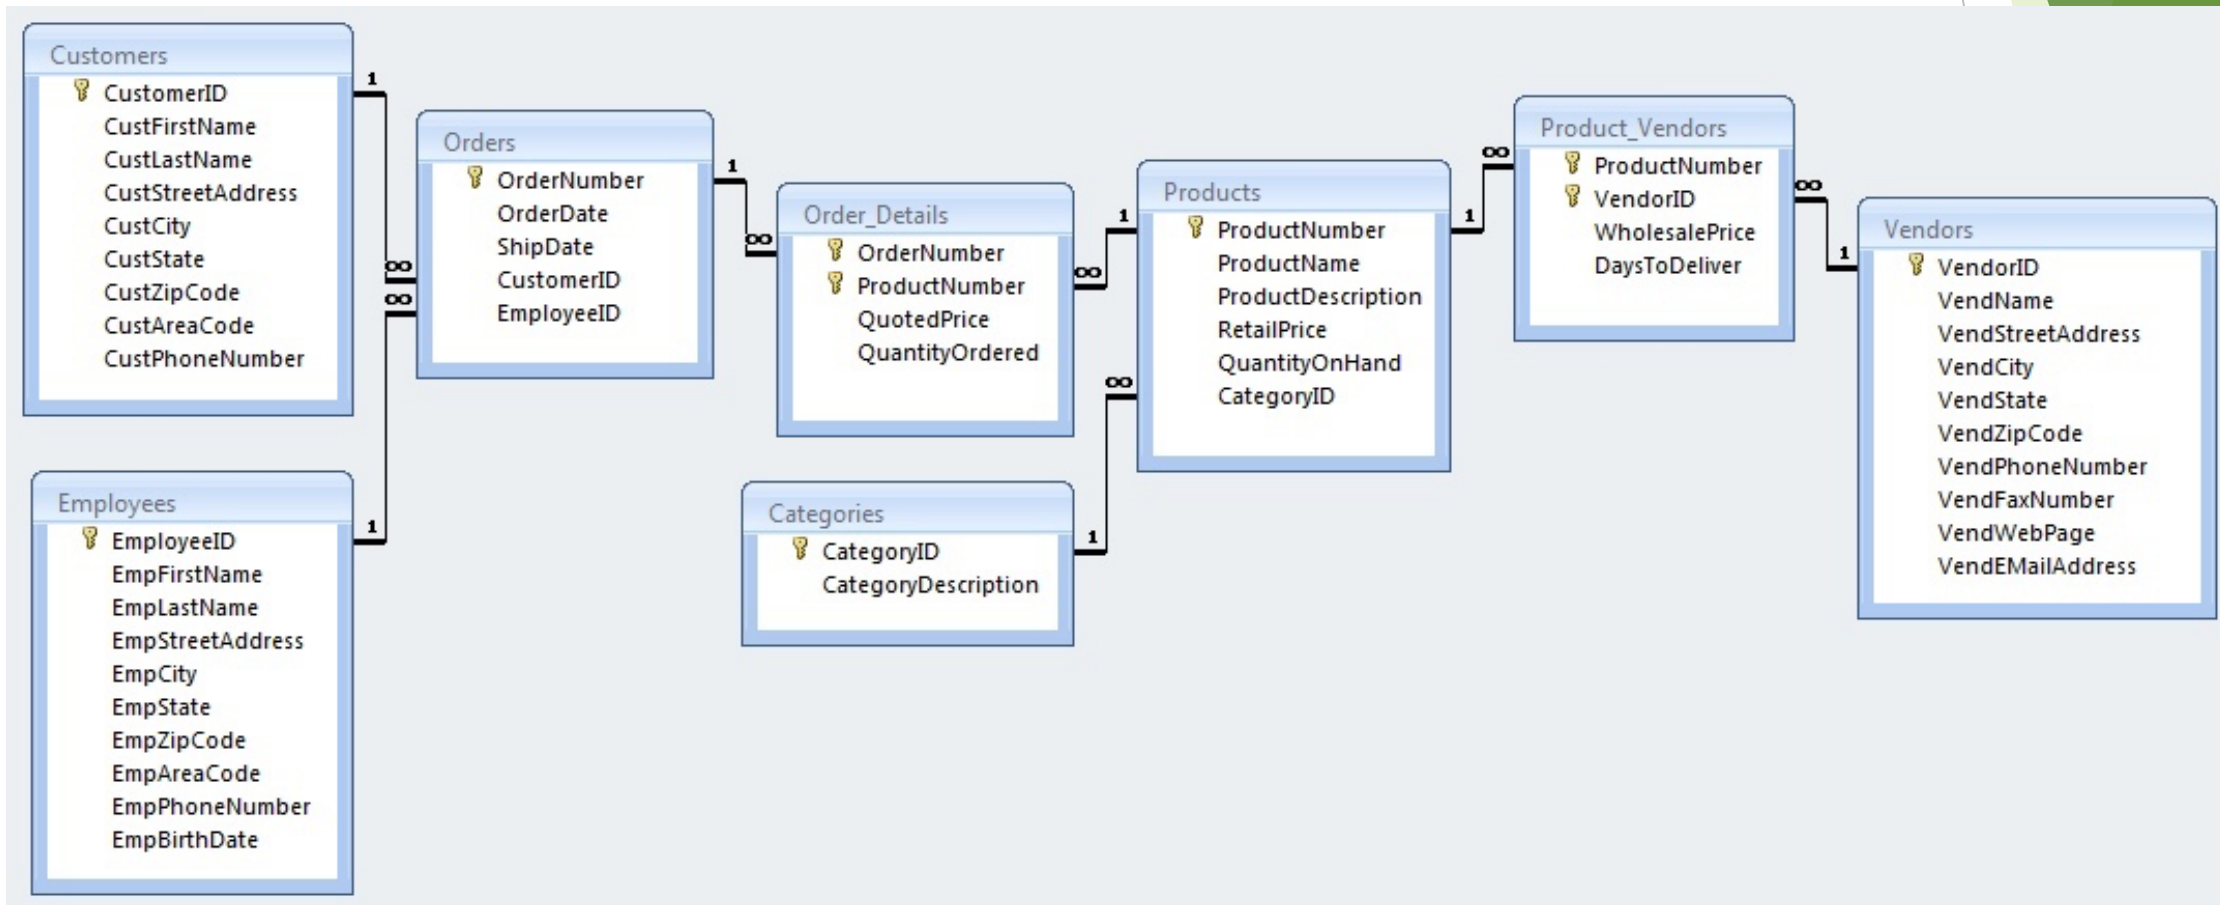
```
SELECT * FROM Orders JOIN Order_Details  
ON Orders.OrderNumber=Order_Details.OrderNumber;
```

Very similar to:

```
SELECT * FROM Orders NATURAL JOIN  
Order_Details;
```


# Designing your data model NATURAL-ly

- Consistent column naming allows you to use NATURAL JOINS.



# CROSS JOIN is like the cartesian product of two sets

- ▶ Take every element (row) of the first set (table) and combine it with every element of the second set.
- ▶ If first set has N elements and second set has M elements, then cartesian product has  $N \cdot M$  elements.
- ▶ There is no “ON” expression to limit results:
  - ▶ `SELECT Orders`  
`CROSS JOIN Order_Details;`



Cartesian Product of Two Sets.

# ON functions are exactly like WHERE

These two expressions are actually equivalent:

- ▶ **SELECT** \* **FROM** Orders  
**JOIN** Order\_Details  
**ON** Orders.OrderNumber=Order\_details.OrderNumber;
- ▶ **SELECT** \* **FROM** Orders  
**CROSS JOIN** Order\_Details  
**WHERE** Orders.OrderNumber=Order\_details.OrderNumber;
- ▶ However, using ON may be more efficient because it tells the DBMS to avoid building the full N·M cartesian product, and just match rows according to a rule.
- ▶ It also makes the join easier to think about, by separating the filtering and JOINing predicates.


# Different JOINS

## SQL JOINS


- ▶ INNER JOIN constructs a table of all pairs of matching rows from two tables.

- ▶ INNER is the default.
- ▶ Useful for *foreign keys* (numeric identifiers)


- ▶ However, there are many other ways to JOIN tables if you don't require matching.




```
SELECT <select_list>  
FROM TableA A  
LEFT JOIN TableB B  
ON A.Key = B.Key
```




```
SELECT <select_list>  
FROM TableA A  
LEFT JOIN TableB B  
ON A.Key = B.Key  
WHERE B.Key IS NULL
```




```
SELECT <select_list>  
FROM TableA A  
INNER JOIN TableB B  
ON A.Key = B.Key
```




```
SELECT <select_list>  
FROM TableA A  
FULL OUTER JOIN TableB B  
ON A.Key = B.Key
```



```
SELECT <select_list>  
FROM TableA A  
RIGHT JOIN TableB B  
ON A.Key = B.Key
```



```
SELECT <select_list>  
FROM TableA A  
RIGHT JOIN TableB B  
ON A.Key = B.Key  
WHERE A.Key IS NULL
```




```
SELECT <select_list>  
FROM TableA A  
FULL OUTER JOIN TableB B  
ON A.Key = B.Key  
WHERE A.Key IS NULL  
OR B.Key IS NULL
```

# LEFT JOIN

- ▶ LEFT JOIN includes **all** rows in the first table (*left*-hand side) and just the matching rows in the second table (right-hand side).

LEFT JOIN


All rows  
from  
First table



Matching rows  
from Second  
table

```
SELECT <select_list>  
FROM TableA A  
LEFT JOIN TableB B  
ON A.Key = B.Key
```

(standard)  
INNER JOIN



```
SELECT <select_list>  
FROM TableA A  
INNER JOIN TableB B  
ON A.Key = B.Key
```



# LEFT JOIN output

- ▶ Like all JOINS, LEFT JOIN prints columns from the left table followed by columns from the right table.
- ▶ However, with LEFT JOIN, some rows will have *NULL* values in the right table columns, meaning that no match was found in the right table.
- ▶ When to use LEFT JOIN?
  - ▶ To supplement a table with additional information that may be available for some rows, **but not available for all the rows.**

staff			
<i>id</i>	<i>name</i>	<i>room</i>	<i>departmentId</i>
11	Bob	100	1
20	Betsy	100	<i>NULL</i>
21	Fran	101	1
22	Frank	102	99999
35	Sarah	200	5
40	Sam	10	7
54	Pat	102	2

department		
<i>id</i>	<i>name</i>	<i>buildingId</i>
1	Industrial Eng.	1
2	Computer Sci.	2
5	Physics	4
7	Materials Sci.	5


- Betsy and Frank have NULLs in the right half of the output because no matching department was found.
- In other words no pair of rows was found to satisfy the ON staff.departmentId=department.id

SELECT \* FROM staff **LEFT JOIN** department ON staff.departmentId=department.id;


staff. <i>i</i> <i>d</i>	staff. <i>nam</i> <i>e</i>	staff. <i>roo</i> <i>m</i>	staff. <i>departmentId</i>	department. <i>id</i>	department. <i>name</i>	department. <i>buildingI</i> <i>d</i>
11	Bob	100	1	1	Industrial Eng.	1
20	Betsy	100	<i>NULL</i>	<i>NULL</i>	<i>NULL</i>	<i>NULL</i>
21	Fran	101	1	1	Industrial Eng.	1
22	Frank	102	99999	<i>NULL</i>	<i>NULL</i>	<i>NULL</i>
35	Sarah	200	5	5	Physics	4
40	Sam	10	7	7	Materials Sci.	5
54	Pat	102	2	2	Computer Sci.	2

# RIGHT JOIN is symmetrical to LEFT

- ▶ Includes all rows from right table and matching rows from left table
- ▶ Reordering the tables makes a RIGHT JOIN a LEFT JOIN, so it is not necessary to use the RIGHT JOIN syntax.



LEFT JOIN



# LEFT JOIN with GROUP BY


In ClassScheduling.slite, count the classes taught by each faculty member:

- ▶ If you want this report to include faculty members teaching zero classes, you must use LEFT JOIN:

```
SELECT StaffID,  
       COUNT(ClassID) AS num_classes  
FROM Faculty NATURAL LEFT JOIN  
Faculty_Classes  
GROUP BY StaffID;
```

- ▶ Note that "COUNT(\*)" would return "1" for faculty members with no classes, because there would still be one unmatched row from the left table.

# LEFT JOIN with exclusion




```
SELECT <select_list>
FROM TableA A
LEFT JOIN TableB B
ON A.Key = B.Key
WHERE B.Key IS NULL
```


- ▶ Includes rows from a table that **must not** match another table.
- ▶ Useful for finding rows lacking something.
- ▶ Just add a **WHERE** clause to look for **NULL** values in the right-hand side of the joined table
- ▶ For example, to determine which faculty members should be assigned a class:
  - ▶ `SELECT * FROM Faculty NATURAL LEFT JOIN Faculty_Classes WHERE ClassID IS NULL;`
- ▶ Which classrooms are unused?
  - ▶ `SELECT * FROM Class_Rooms NATURAL LEFT JOIN Classes WHERE ClassID IS NULL;`

# FULL OUTER JOINS are not available in MySQL or SQLite


- You can *emulate* FULL OUTER JOIN with the UNION of two queries.




```
SELECT <select_list>  
FROM TableA A  
LEFT JOIN TableB B  
ON A.Key = B.Key
```




```
SELECT <select_list>  
FROM TableA A  
LEFT JOIN TableB B  
ON A.Key = B.Key  
WHERE B.Key IS NULL
```




```
SELECT <select_list>  
FROM TableA A  
INNER JOIN TableB B  
ON A.Key = B.Key
```




```
SELECT <select_list>  
FROM TableA A  
RIGHT JOIN TableB B  
ON A.Key = B.Key
```



```
SELECT <select_list>  
FROM TableA A  
RIGHT JOIN TableB B  
ON A.Key = B.Key  
WHERE A.Key IS NULL
```



```
SELECT <select_list>  
FROM TableA A  
FULL OUTER JOIN TableB B  
ON A.Key = B.Key
```



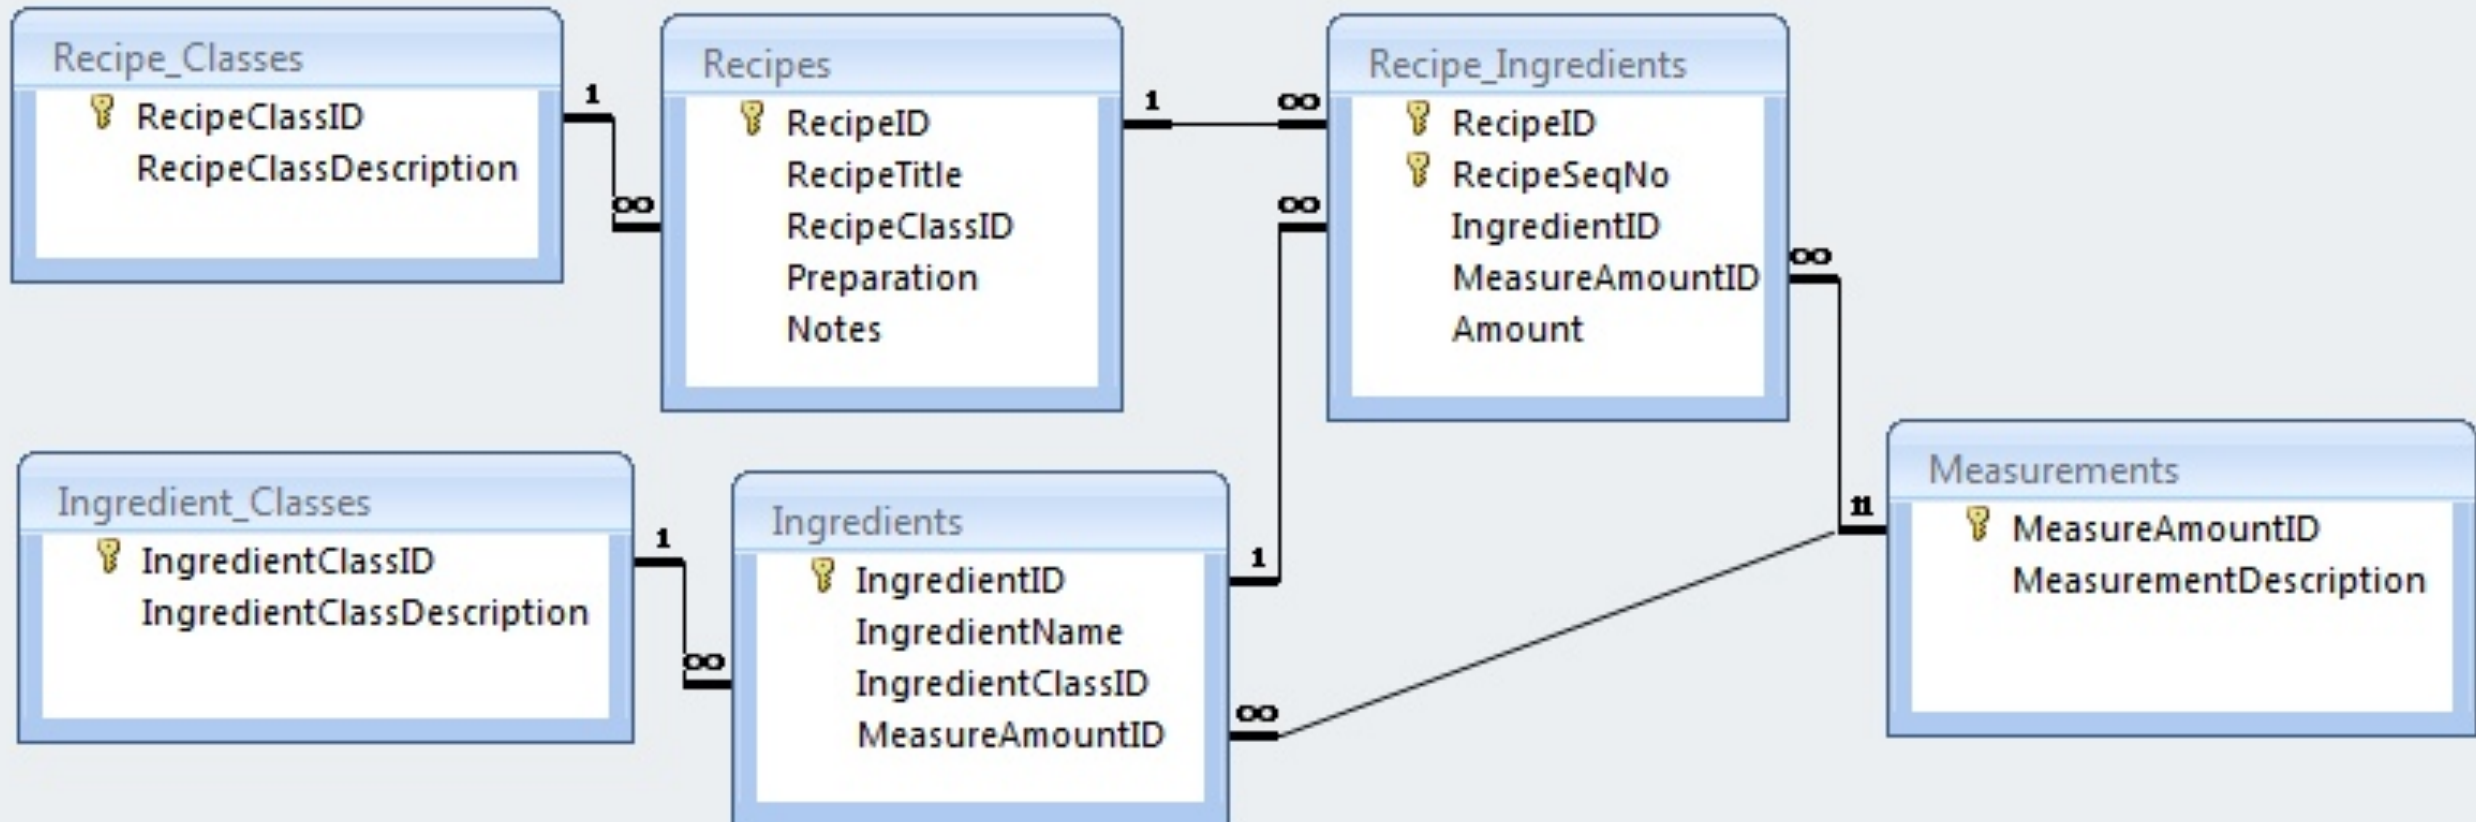
```
SELECT <select_list>  
FROM TableA A  
FULL OUTER JOIN TableB B  
ON A.Key = B.Key  
WHERE A.Key IS NULL  
OR B.Key IS NULL
```

# Examples

The background features abstract, overlapping green geometric shapes, primarily triangles and polygons, in various shades of green, ranging from light lime to dark forest green. These shapes are concentrated on the right side of the slide, with some extending towards the left. The overall effect is a modern, minimalist design.

# Recipes.sqlite: List the number of recipes in each category (RecipeClassID)

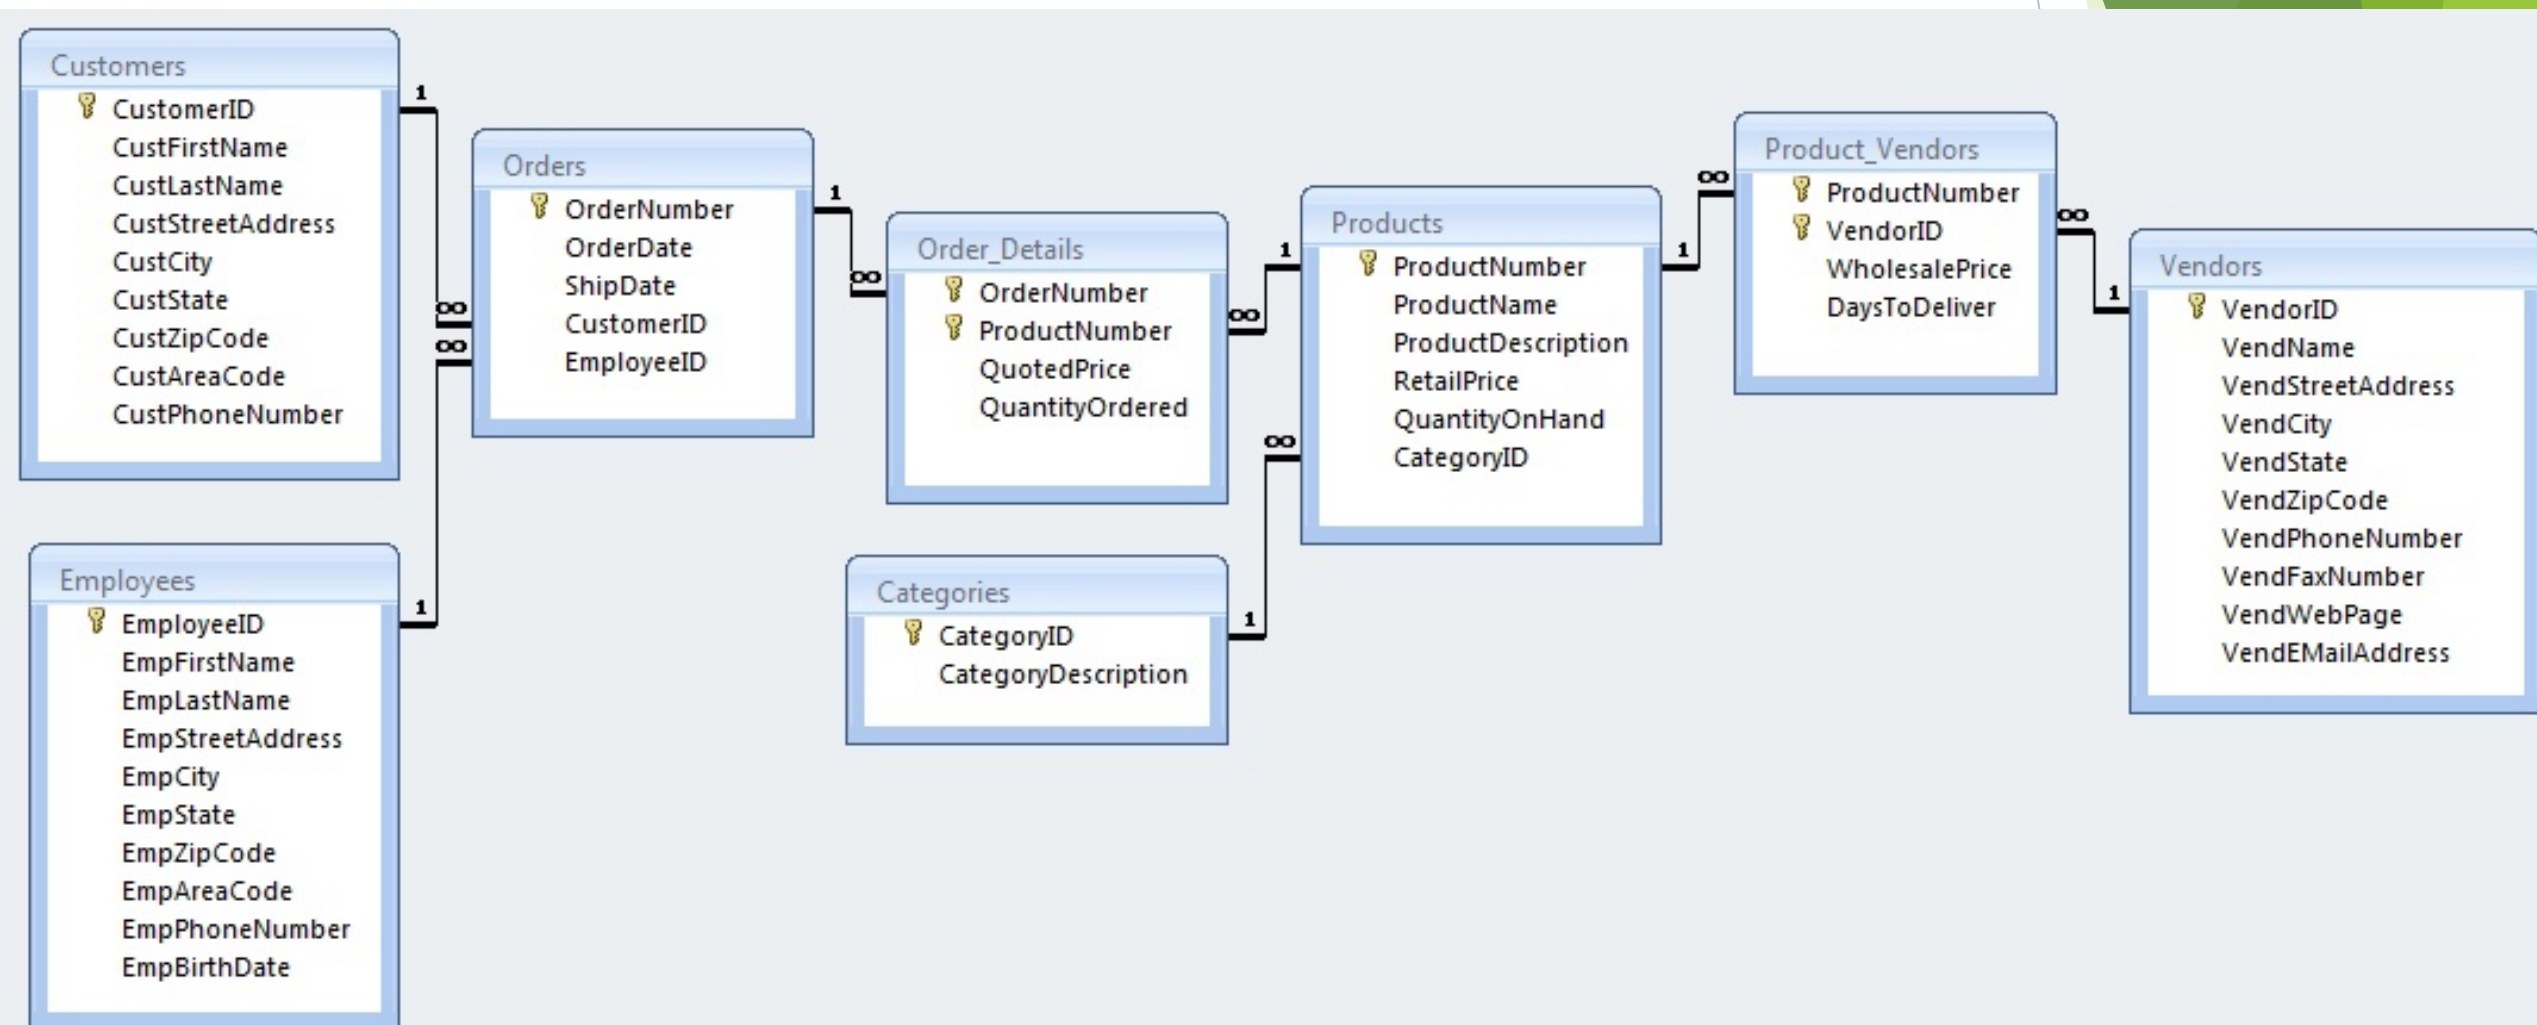
```
SELECT RecipeClassDescription, COUNT(RecipeID) AS RecipeCount  
FROM Recipe_Classes LEFT NATURAL JOIN Recipes GROUP BY RecipeClassID;
```





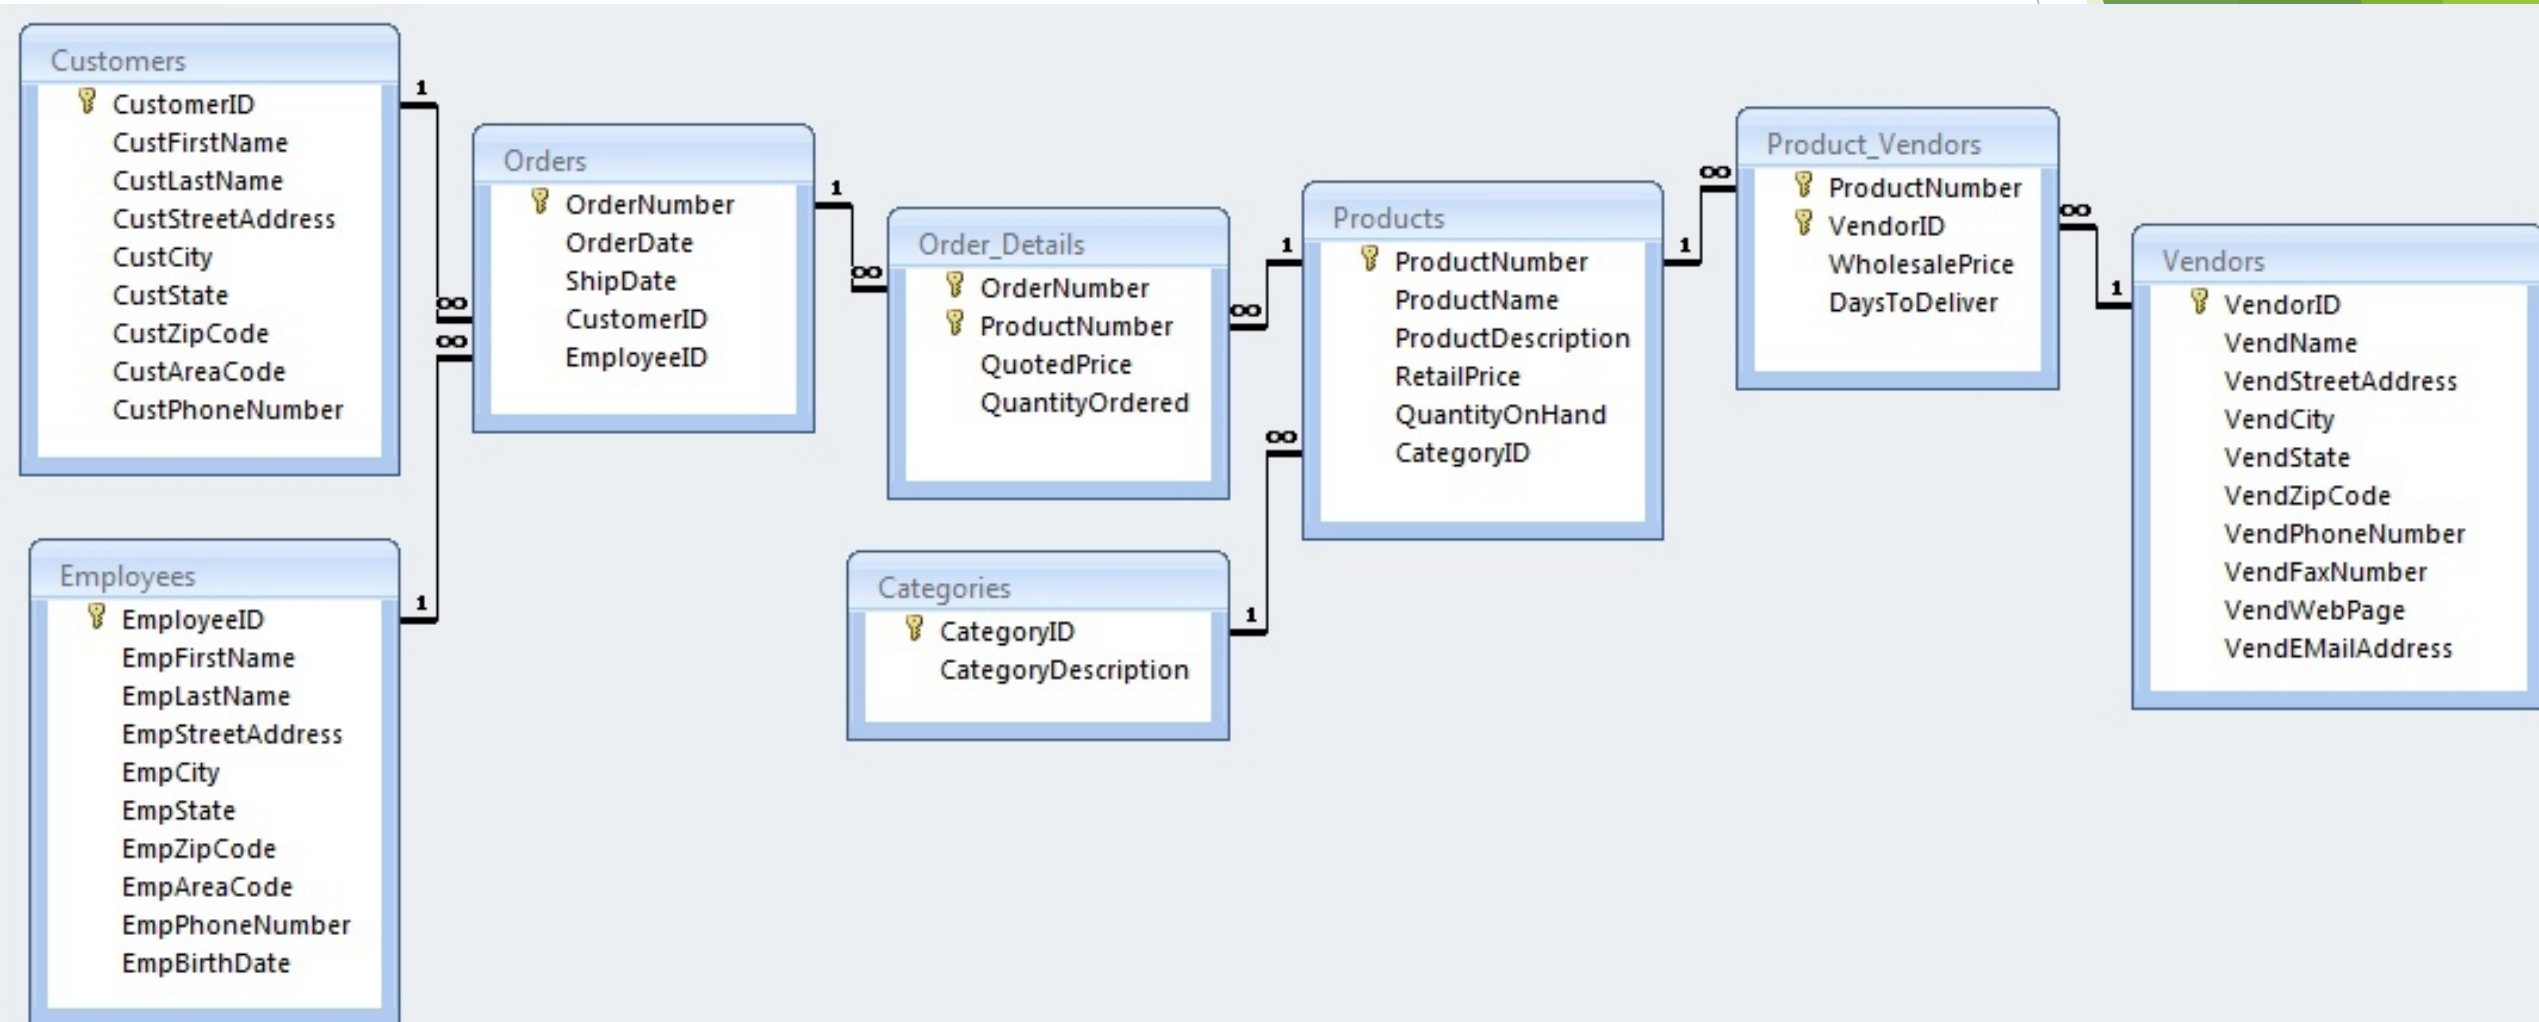
# SalesOrders.sqlite: For all products, list any orders of that product and their dates.

```
SELECT Products.ProductNumber, ProductName, OrderDate FROM  
Products LEFT NATURAL JOIN (Order_Details NATURAL JOIN Orders);
```



# Display customers who have no sales rep (employees) in the same ZIP Code.


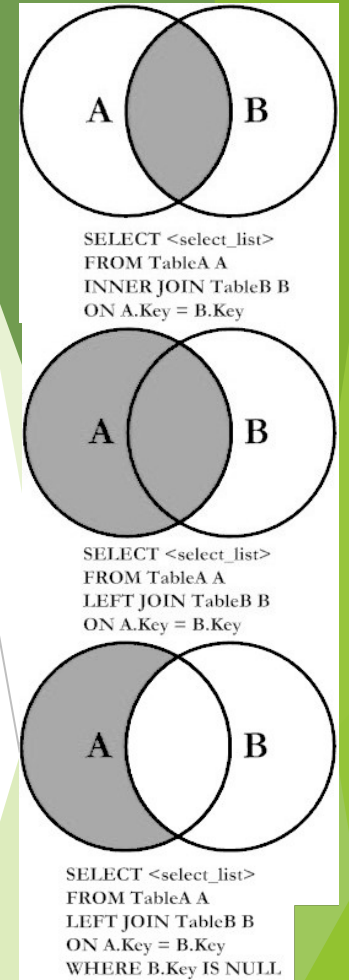
```
SELECT * FROM Customers LEFT JOIN Employees ON  
CustZipCode=EmpZipCode WHERE EmpZipCode IS NULL;
```



# Recap

Introduced different types of JOINS:

- ▶ **INNER (default)**: prints all pairs of rows (one from first table, one from second table) that satisfy the *JOIN predicate*.
- ▶ **LEFT**: same as INNER, but adds rows from LEFT table that never satisfied the JOIN predicate.
- ▶ **LEFT with exclusion**: only print rows from left table that never satisfied the JOIN predicate.
- ▶ **CROSS JOIN**: print the cartesian product, meaning all rows from the first table combined with all rows from the second table. There is no “ON” to match rows.



Cartesian Product of Two Sets.